

## Project Development Phase

### Model Building and Training

Date	19 February 2026
Team ID	LTVIP2026TMIDS77295
Project Name	Smart Sorting: Transfer Learning for Identifying Rotten Fruits and Vegetables.

## Model Building

### 1. Choice of Architecture

The model is built using transfer learning on top of a pre-trained MobileNetV2 convolutional neural network. MobileNetV2 is chosen because it is lightweight, optimized for image classification, and already trained on a large dataset (ImageNet), so it has learned rich low-level features like edges, textures, and shapes that transfer well to fruit images.

- The base MobileNetV2 is loaded with pre-trained weights and without its original classification head.
- The base layers are initially frozen, so their weights are not changed during the first training phase. This helps avoid destroying useful pre-learned features and speeds up convergence on a small to medium-sized dataset.

### 2. Custom Classification Head

A custom classification head is added on top of the MobileNetV2 base to adapt it to the binary classification problem (Fresh vs Rotten):

- Global average pooling layer to reduce the spatial feature maps into a single vector.
- One or more dense (fully connected) layers with ReLU activation to learn task-specific patterns.
- Dropout (optional) to reduce overfitting by randomly deactivating neurons during training.
- Final dense layer with 1 neuron and sigmoid activation (or 2 neurons with softmax) to output the probability of the image being fresh or rotten.

The entire model is compiled with:

- Loss function: Binary cross-entropy (for two classes).
- Optimizer: Adam optimizer with an appropriate learning rate (e.g., 1e-4).
- Metrics: Accuracy as the primary metric, with optional precision, recall, or F1 score in analysis.

### 3. Input Preprocessing

Before feeding images into the model, they are preprocessed to match MobileNetV2 requirements:

- Images are resized to a fixed resolution (e.g., 224×224).
- Converted to RGB and normalized (pixel values scaled to a suitable range, such as [0, 1] or the specific preprocessing function provided by MobileNetV2).

- Data is loaded using generators (like `ImageDataGenerator.flow_from_directory`) that read images from train and test/validation directories and apply augmentation in real time.
- 

## Model Training

### 1. Training Configuration

The model is trained on the cleaned and split dataset:

- Training images in dataset/train/fresh and dataset/train/rotten.
- Validation or test images in dataset/test/fresh and dataset/test/rotten.
- Typical hyperparameters:
  - Batch size: e.g., 16 or 32.
  - Epochs: e.g., 10–30 for initial training, adjusted based on convergence.
  - Learning rate: small value (e.g., 1e-4) to ensure stable training.

The training loop uses:

- `model.fit()` with training and validation generators.
- Callbacks like `ModelCheckpoint` (to save best weights) and `EarlyStopping` (to stop when validation loss stops improving) for better generalization.

### 2. Data Augmentation During Training

To improve robustness and performance, data augmentation is applied on the fly to training images:

- Random rotations, horizontal flips, zoom, width/height shifts, and brightness changes.
- This artificially increases data variety and helps the model generalize to different real-world conditions (angles, lighting, background).
- Augmentation is applied only to the training generator, not to validation/test data.

### 3. Initial Training Phase (Frozen Base)

In the first phase:

- All layers of the MobileNetV2 base are frozen; only the newly added classification head is trainable.
- The model is trained for several epochs until training and validation accuracy stabilize.
- This phase lets the top layers learn to map MobileNetV2 features to the fresh/rotten decision without disturbing the pre-trained weights.

### 4. Fine-Tuning Phase (Optional, for Higher Performance)

For improved performance, a second fine-tuning phase can be carried out:

- Unfreeze the upper few layers of the MobileNetV2 base while keeping lower layers (which capture generic features) frozen.
- Re-compile the model with a lower learning rate (e.g., 1e-5) to prevent large, destabilizing weight updates.
- Train for additional epochs while monitoring validation loss and accuracy.

This fine-tuning allows the model to adapt deeper feature representations specifically to the fruit freshness detection task and can significantly boost validation accuracy.

## 5. Saving the Trained Model

After training:

- The best performing weights (based on validation accuracy/loss) are saved to a file such as `model/healthy_vs_rotten.h5`.
  - This file is later loaded by the Flask backend (`predict.py`) for real-time predictions without retraining.
- 

## Performance Evaluation (Under Performance Testing)

Once training is complete, performance testing focuses on how well the model generalizes:

- Accuracy on the validation/test set: fraction of correctly classified images.
- Confusion matrix: counts of true positives, true negatives, false positives, and false negatives for Fresh and Rotten, revealing where the model makes mistakes.
- Classification report: precision, recall, and F1 score per class to understand performance especially on the minority class.

These metrics are used to:

- Decide whether the model is good enough for deployment.
- Analyze failure cases (e.g., rotten fruits misclassified as fresh) and guide dataset improvements or further fine-tuning.